

Project Proposal

AgentDesign: Hermes Profile Architecture

Math & Librarian Profiles — Design Documenta

Mehdi Ghasemi

Project Proposal

September 2026

Executive Summary

This document records the complete design of the two Hermes agent profiles that power my mathematical research workflow — `math` and `librarian` — and specifies an abstraction layer under which either profile can be re-pointed at any future math project without modification. Sanitized, ready-to-use copies of both profiles ship with this document as the `math/` and `librarian/` folders (Section 4); the companion artifacts listed in Section 11 let every claim here be re-derived from a live installation.

The architecture resolves into five layers: a shared *Hermes core*, an isolated *profile system*, a *skill catalog* (capability-gated per profile), an *MCP bridge* (native servers plus a generic CLI-adapter pattern), and an *external service plane*. The math profile is a full-stack researcher with a sandboxed toolchain environment; the librarian profile is a capability-degraded derivative that shares every registration but disables the computational surface and adds ingestion skills of its own.

Two structural findings drive the abstraction design: (i) the librarian profile's MCP servers import scripts from the math profile's directory tree, so it cannot exist without the math profile; and (ii) both profiles carry a stale reference to a development worktree that no longer exists (a merge left the registered path dangling). Section 10 proposes a capability/manifest model — persona templates, path interpolation, a shared adapter, a service registry, and per-project binding files — that removes every identified coupling in five concrete steps.

Contents

1 Objectives and Scope	3
2 Layered Architecture	3
3 Profile System: Structure Comparison	4
4 Folder Structure	5
5 Skill System: Design for General Math Projects	7
6 MCP Bridge	9
7 Tool Stack and Service Plane	10
8 The Research Workflow	11
9 Identified Couplings and Defects	13
10 Toward an Abstraction Layer	13
11 Roadmap for AgentDesign Itself	16

I Objectives and Scope

The `AgentDesign` project has three objectives:

1. **Document the as-is architecture.** Record, with file-level evidence, how the math and librarian profiles are composed: directory layout, configuration sections, skill catalogs, MCP server registrations, external services, scheduled jobs, and state stores. This is Sections 2–8.
2. **Identify the coupling points.** Enumerate every place where a profile depends on the specific current project layout — absolute paths, cross-profile imports, hardcoded endpoints, and stale references — so that the design of record distinguishes what is load-bearing from what is incidental (Section 9).
3. **Specify an abstraction layer.** Define how profiles become reusable personas: independent of any particular math project, instantiable for a new project with zero or one line of configuration, and testable against a capability contract (Section 10).

The scope is the two active research profiles. The remaining four registered profiles (`default`, `coder`, `advisor`, `trader`) are out of scope; their existence is noted only in Section 3. This project does not modify any live profile configuration — it is a documentation and design artifact.

2 Layered Architecture

Figure 1 gives the visual layout of the infrastructure under which both profiles operate. Five layers are distinguished; each lower layer is shared, each upper layer is isolated per profile at `~/.hermes/profiles/<name>/`.




Figure 1: Five-layer profile architecture. Layers 2–3 are the capability surface of a profile; layers 0 and 4 are shared infrastructure. Arrows indicate the direction in which each layer is configured by, or depends on, the layer above it.

Layer 0 — Hermes Core

The core lives at `~/.hermes/hermes-agent/`: agent runtime, CLI (profile switching, tool registry), platform gateway (Telegram/WhatsApp/webhook adapters), built-in tools (terminal, browser, vision, file I/O, web

search, kanban, memory) and the MCP discovery layer (`tools/mcp_tool_*.py`). Shared binaries (`tirth` security scanner, `uv`, Python 3.14, Node 26, `ripgrep`, Chromium) live in `~/.hermes/tools/`; all MCP stdio servers run from a single shared virtualenv at `~/.hermes/mcp-servers/.venv/`.

Layer 1 — Profile System

Each profile is an isolated directory with its own configuration, credentials, state and sandbox. The two profiles under study:

- `math` — full-stack mathematical researcher. It owns a sandboxed home environment carrying the local toolchain (Section 3), an active messaging gateway, and the scheduled literature jobs.
- `librarian` — literature manager (“manage and maintain books and articles”). It keeps an empty sandbox home, registers no scheduled jobs and no messaging platform, and reuses the same kanban assignment as `math` — itself a coupling noted in Section 9.

Both profiles run the same local model through LM Studio, with an `on_session_start` hook that guarantees the model server is up before each session. Long-term memory in both profiles is backed by Honcho (`memory.provider: honcho`).

Layer 2 — Skill Catalog

Skills are Markdown-defined capability units (`SKILL.md` with YAML frontmatter), organized in category directories, optionally carrying scripts and credential files. Each profile holds the same catalog — nine populated categories plus a handful of top-level skills, and the same empty category shells — and differs only in which entries are activated. The per-profile inventory extracted from a live installation is kept in `references/skill_inventory.md`; Section 5 summarizes the design-relevant parts.

Layer 3 — MCP Bridge

Skills reach services either directly (terminal/shell) or through Model Context Protocol servers registered in each profile’s config; every registered tool is exposed to the model under a name built from the server and the tool, joined by double underscores — for example, the `vikunja` task-creation tool. Section 6 documents the two server families and the cross-profile dependency this layer creates.

Layer 4 — Service Plane

A fixed set of self-hosted services on the local network, plus SaaS APIs and a local inference server. Table 4 lists them with their role in the workflow.

3 Profile System: Structure Comparison

Table 1 compares the two profiles as designs rather than as installations: the question is what each profile *is*, not what it happens to contain today.

Table 1: Profile design comparison: the librarian is a capability-degraded derivative of the math profile, not an independent design.

Design element	math	librarian
Role	Full-stack mathematical researcher	Literature manager: “manage and maintain books and articles”
Inference	Local model served by LM Studio	Same model and provider, with an explicit API mode
Skill surface	Full catalog; computation, formal proof and manuscript skills active	Same catalog with the computational surface disabled in configuration; adds its own ingestion skills (<code>calibre-library-ingestion</code> , <code>djvu-to-pdf</code>)
MCP bridge	Native servers plus adapter-wrapped CLI tools	Identical registrations, all resolving into the math profile’s tree
Plugins	<code>vikunja</code> and <code>wiki-browser</code> enabled	Identical
Toolchain sandbox	Lean 4, Rust, SageMath and TeX Live reachable from the profile home	Empty stub
Scheduled jobs	Literature scan and wiki maintenance	None
Messaging platform	Active gateway	None
Cross-session memory	Honcho	Honcho
Kanban assignment	math	math — inherited, not chosen

The librarian is a capability-degraded derivative of math, not an independent design: it inherits every skill directory and MCP registration, then suppresses the computational surface via the `skills.disabled` list. It also carries two skills that exist only in its own tree — `calibre-library-ingestion` and `djvu-to-pdf` (productivity category) — which is where its genuine specialization lives.

The math profile carries a sandboxed home directory that stands in for a machine-local toolchain. It holds a Lean 4 installation managed by `elan`, a Rust toolchain together with its registry and download caches, a dedicated browser profile for NotebookLM grounded synthesis, runtime state directories for SageMath and Maxima, and the font and format caches of TeX Live. Two properties matter for the abstraction design. First, most of the space is disposable package cache rather than toolchain state: the load-bearing parts are the Lean and Rust installations, everything else can be re-provisioned. Second, the sandbox is not self-contained — the SageMath and TeX Live binaries resolve to host-level installations rather than to anything inside the profile home — so a profile’s computational capability depends on the host even when it looks self-contained.

4 Folder Structure

Three directory levels carry the design: the shared Hermes root, the two profile trees, and the per-project binding file. The listings below show the structure with the design-relevant entries; the profile skeleton in Listing 4 marks which entries exist in which profile.

Shared root

The `~/hermes/` root holds everything that is not profile-scoped. Two entries are load-bearing for the coupling analysis: `mcp-servers/` holds the single Python environment that every MCP server of every profile runs in,

and `skills/` is the shared pool that both profiles fall back to for both skills and plugins.

```

1 ~/.hermes/           # shared Hermes root
2  config.yaml         # global defaults; profiles override
3  profiles/          # one directory per profile (math, librarian, advisor, ...)
4  skills/             # shared skill pool: 17 categories
5  plugins/            # shared plugin installs: vikunja, wiki-browser, hermes-
    achievements
6  desktop-plugins/    # desktop plugin surfaces: vikunja, wiki-browser
7  plugin-data/        # per-plugin persistent state: hermes-achievements
8  mcp-servers/.venv/  # ONE Python env for all 13 MCP servers
9  memories/ sessions/ state.db # memory files, session store, global state
10 kanban/ kanban.db   # kanban board + dispatcher state
11 cron/ hooks/ agent-hooks/ # scheduling and lifecycle hooks
12 bin/ node/ webui/ tools/ # runtime binaries, desktop build, CLI tools

```

Profile tree: the two profiles side by side

The math profile is a structural superset: it carries the runtime residue of an active profile, while the librarian keeps only what its role needs. Two differences are design-relevant — `scripts/`, the MCP adapter directory, exists only on the math side, and `home/` is populated only there. Those two facts are the structural root of the path coupling in Section 6.

1	profiles/math/		profiles/librarian/
2	config.yaml	profile.yaml	config.yaml profile.yaml
3	.env		.env
4	skills/	24 top-level dirs	skills/ 23 top-level dirs
5	plugins/	vikunja, wiki-browser,	plugins/ empty (borrows shared pool)
6		hermes-android (disabled)	
7	desktop-plugins/	vikunja, wiki-browser	desktop-plugins/ [absent]
8	scripts/	mcp_cli_adapter.py + 8 .py	scripts/ [absent]
9	home/	toolchain sandbox	home/ (empty stub)
10	scratchpad/	failure scratchpad	scratchpad/ [absent]
11	cron/ hooks/ memories/ sandboxes/		cron/ hooks/ memories/ sandboxes/
12	sessions/ workspace/ cache/ logs/		sessions/ workspace/ cache/ logs/
13	backups/ state.db SOUL.md		backups/ state.db SOUL.md
14	telemetry/	[absent]	telemetry/ (sender on)
15	vault/	[absent]	vault/ (credential store)
16	state.db / SOUL.md / backups/		state.db / SOUL.md / backups/

Beyond the entries shown, the librarian keeps `telemetry/` and `vault/`; math's extra entries are dominated by runtime residue (gateway sockets, backup snapshots, retired write-ahead logs) plus the three design-relevant directories listed above.

Skill tree layout

Both profiles carry the same 24-directory layout except for the single extra top-level skill in math; activation, not layout, is what differs.

```

1 profiles/<profile>/skills/ # 24 top-level dirs (math) / 23 (librarian)
2  research/ productivity/ mlops/ # 9 populated categories; each holds
3  software-development/ github/ devops/ # <skill>/SKILL.md with optional
4  web/ autonomous-ai-agents/ social-media/ # scripts/ and references/
5  apple/ creative/ data-science/ email/ # 9 empty category shells (reserved,
6  environment/ media/ note-taking/ # no SKILL.md; used by the curator

```

```

7  mathematical-research/ smart-home/          # as retirement targets)
8  differential-sdp/ ghasemi-latex-style/       # 6 top-level skills in math,
9  hermes-plugin-management/ semantic-scholar/# 5 in librarian
10 workspace-doc-hierarchy/ irene-rewrite-dev/ [math only]
11 .archive/                                     # retired skills kept as curator tombstones
12 .curator_ledger.jsonl .curator_state/ .curator_suppressed/
13 .bundled_manifest .usage.json .hub/ .locks/

```

Distributable profile packs

The `math/` and `librarian/` folders inside this project are sanitized, ready-to-use copies of the two profiles: the same directory structure, the same configuration and skill trees, with every credential, endpoint, hostname and local path replaced by a placeholder. A builder script in `references/` copies the configuration, skill, plugin, script and memory trees from a live installation, scrubs secrets and identity, and writes a per-pack `README.md`. Copying a pack into `~/.hermes/profiles/` and filling in the placeholders reproduces the profile; the packs are described in Section 4.

This document's own folder

```

1 AgentDesign/
2  agent_design.tex          # this document (source of record)
3  agent_design.pdf         # compiled output
4  AGENTS.md                # project guidance for future sessions
5  math/ librarian/         # sanitized, ready-to-use profile packs (Section~\ref{sec:
    packs})
6  stack/                   # the profiles' service plane as one compose project
7  cpistuff/                # CPI LaTeX template: cpi.sty + logos
8  t1ggm.fd ts1ggm.fd       # Garamond T1 declarations (compile prerequisite)
9  data/profile_comparison.json # machine-readable extraction (secret values redacted)
10 references/extract_profile_data.py references/build_profile_packs.py
11 references/sanitize_extracted_data.py
12 references/skill_inventory.md # per-profile skill inventory
13 .gitignore

```

The `stack/` folder turns the service plane of Section 7 into an executable artifact: one compose project that brings up every container the two profiles talk to, with credentials interpolated from an untracked `.env`. Its `README.md` maps each container to the MCP server or skill that consumes it and records what the live deployment gets wrong.

Reproducing the build requires one environment step beyond the template: the CPI style loads the Garamond T1 fonts, which are not part of a stock TeX Live installation. They are installed per-user (no root) under `~/texmf/` — `ggm*.tfm` in `fonts/tfm/ggm/`, `ggm*.vf` in `fonts/vf/ggm/`, `ggm*.pfb` in `fonts/type1/ggm/`, `ggm.map` in `dvips/ggm/`, then `mktexlsr ~/texmf`. Without it the build fails at the first page with a missing-font error rather than a document error.

5 Skill System: Design for General Math Projects

Skills are the unit of reusability. Each `SKILL.md` carries a trigger description in its frontmatter, an imperative body, optional scripts, and — where needed — a credential file (`.env`) co-located with the skill that consumes it. The catalog is shared between the two profiles: the same populated categories, the same empty category shells, and the same top-level skills, with a single exception — `irene-rewrite-dev` is an active top-level skill

in math and a curator tombstone in the librarian. Only activation differs. Table 2 summarizes the catalog by design role; the per-entry inventory extracted from a live installation is in `references/skill_inventory.md`.

Table 2: Skill catalog by category, with activation per profile.

Category	Design role	Activation
research	The research surface proper: the six-stage workflow, symbolic and numeric computation, formal verification, literature retrieval and the wiki pipeline	full in math; computational and proof skills disabled in librarian
productivity	Knowledge management: task ledger, notes, bibliography, manuscript production, document conversion	full in both; the librarian adds its ingestion skills
mlops	Local inference operations and pipeline tooling	reduced in librarian
software-development	Hermes self-extensibility, planning, subagent orchestration, desktop inspection	full in math; partly disabled in librarian
github, devops, agents, web	Repository operations, SDLC review, delegation, blocked-page recovery	mixed; review and delegation off in librarian
social-media	Community reading	inverted gate: disabled for math, enabled for librarian
top-level skills	Cross-cutting capabilities: differential-algebraic SDP, manuscript style, plugin management, citation graph, document hierarchy	full in math; the librarian disables the development-oriented entry

Three observations are design-relevant:

1. **The math research surface is exactly the set of skills a generic math project needs:** decompose → literature → compute → verify → write, implemented by `math-research-workflow`, whose frontmatter names the sibling skills it relies on. This set, not the full catalog, defines the reusable “mathematician” capability contract (Section 10).
2. **The librarian’s disabled list is a negative specification of that same contract:** every computational skill (`sagemath-mcp`, `sympy-mcp`, `lean4`, `lean4-goedel-agent`, `scientific-coding`), the DSDP skills, and the Irene/dev skills are turned off. Capability gating by negation works today but is brittle in a subtler way than it looks: entries resolve by *leaf name* against the profile’s skill tree, which includes the curator tombstone directory `skills/.archive/`. Every disabled name therefore resolves to a real `SKILL.md`; none dangles. But a sizeable minority of them resolve *only* under `.archive/` — they are residue from skills the curator has already retired (`docx`, `opencode`, `test-driven-development`, and for the librarian also `irene-rewrite-dev`, which is an *active* top-level skill in math while pointing at a librarian tombstone). Because resolution is name-based, any of these stale entries would silently disable a same-named skill if it were ever reinstalled. One inversion deserves note: `reddit-reading` is disabled for math but enabled for librarian — the only skill whose gate runs opposite to every other difference between the two profiles.
3. **One reference is stale:** `math-research-workflow` names `simplerag-memory` in its frontmatter and body, several times as shell invocations of a client script that does not exist; no such skill exists in either profile tree or in the bundled pool.

6 MCP Bridge

Thirteen MCP servers are registered per profile, all launched from one shared interpreter environment. They fall into three families:

Table 3: MCP server registrations. “Pattern” distinguishes native servers from generic CLI-adapter wraps; target paths are written relative to the profile tree.

Server	Pattern	Target	Tools
sagemath	native	skills/research/sagemath-mcp/	10
lean4	native	skills/research/lean4/	4
searxng	native	skills/research/searxng/scripts/	1
vikunja	CLI adapter	skills/productivity/vikunja/	15
wolfram-alpha	CLI adapter	skills/research/wolfram-alpha/	3
lightrag-query	CLI adapter	skills/research/lightrag/	3
siyuan	CLI adapter	skills/productivity/siyuan/	7
zimi	CLI adapter	skills/research/zimi/	4
zotero	CLI adapter	skills/productivity/zotero/	7
scientific-coding	CLI adapter	skills/research/scientific-coding/ (CPU and memory limits via env)	4
latex-manuscript	CLI adapter	skills/productivity/latex-manuscript/	5
hermes-irene	direct script	inside a development worktree since merged and removed: registered, but dangling	—
calibre	native, librarian only	math: environment block without a command; librarian: a server with its own virtual environment	12

The **generic CLI-adapter pattern** is the load-bearing design of this layer. A single script — the generic adapter kept in the profile’s `scripts/` directory — wraps any argparse-based Python CLI tool as a stdio MCP server: it introspects the target’s parser at startup, exposes each subcommand as an MCP tool with JSON-schema arguments, and reconstructs command lines from tool calls. Most of the registered servers use it; adding a new service-backed skill therefore requires only a thin CLI script plus a config entry (endpoint and token via environment), with no protocol code.

Two cross-profile facts follow directly from the table:

- Every path in every MCP registration — including all of the librarian’s — points into the **math profile’s directory tree**: every adapter tool script, every native server script, and the adapter itself. The librarian has no `scripts/` directory of its own. Deleting or moving the math profile would silently break every tool-bridged capability of the librarian.
- Both profiles register `hermes-irene` against a script inside a development worktree that was merged into the main `Irene` package and then removed; the file was never committed, so the registration dangles. It is dormant in math and left enabled in the librarian, where it fails at startup. The correct target is the merged package, where no such script has yet been created.

7 Tool Stack and Service Plane

Table 4 lists the services both profiles reach, grouped by role in the research workflow. Endpoints and credentials are not part of the design: each profile resolves them from its own environment file at load time, so the same registrations serve any deployment. The profiles share the bulk of their environment keys; the math-only keys are messaging and auxiliary settings (Telegram bot token and allowlist, WhatsApp mode fields, an OpenRouter key).

Table 4: Service plane, grouped by role in the research workflow. Both profiles share the same registrations; differences are noted.

Service	Role in the workflow	Integration
LightRAG	Graph-RAG over the curated literature lake: ingestion and multi-hop query	MCP
SearXNG	Federated web metasearch with structured output	MCP
ZIMI	Offline encyclopedia for archival grounding before any live search	MCP
Vikunja	Task ledger: hypothesis trees, phase tracking, cross-session state	MCP and plugin
SiYuan	Block-level notes: human-readable findings and equation blocks	MCP and plugin
Calibre	Ebook library management and ingestion	librarian only
Zotero	Bibliography, BibTeX export, citation verification	MCP
Wolfram Alpha	Symbolic and numeric computation beyond the local stack	MCP
Honcho	Persistent cross-session research memory	built-in memory provider
LM Studio	Local inference for both profiles	model provider
Portainer	Container management of the service plane itself	terminal
Semantic Scholar, arXiv, Crossref	Citation-graph layer and the metadata verification gate	skills

The **four integration channels** a skill may use, in the preference order codified by `math-research-workflow`:

1. *MCP native / adapter tools* — fastest path into services; no shell involved.
2. *Terminal invocations of CLI scripts* bundled with skills (e.g. the arXiv research script, wiki ingestion pipeline).
3. *Built-in Hermes tools* — web search/extraction, file I/O, browser automation.
4. *Local binaries from the sandbox home* (SageMath, Lean 4 via `elan/rustup`, TeXLive) for computation and compilation that must stay on-machine.

Plugins

Plugins are the third extension channel beside skills and MCP servers, and their configuration is byte-identical in the two profiles.

Table 5: Plugin configuration.

Element	State
Enabled plugins	<code>vikunja</code> , <code>wiki-browser</code> — identical in both profiles
Disabled plugins	<code>hermes-android</code> , <code>spotify</code> — identical in both profiles
Tool-override entries	<code>hermes-android</code> and <code>wiki-browser</code> , both with overrides denied
Profile plugin directory	<code>math</code> carries the plugin trees; the librarian's is empty
Shared installs	resolved from the shared Hermes plugin pool when a profile has no local copy
Desktop surfaces	<code>desktop-plugins/</code> carries the <code>vikunja</code> and <code>wiki-browser</code> surfaces

Three facts matter for the abstraction design. First, the librarian enables both plugins while its own `plugins/` directory is empty: plugin resolution falls back to the shared pool exactly as skill resolution does, so the librarian's plugin capability is entirely borrowed — and, as with skills, the coupling is by name rather than by manifest. Second, `spotify` is listed as disabled in both profiles although no installed copy exists in either profile tree or in the shared pool; the name matches an upstream catalog plugin inside the Hermes source tree, so the entry is a no-op rather than a broken install — the same class of residue as the curator tombstones noted in Section 5. Third, `hermes-android` sits in `math`'s `plugins/` directory but is disabled, while the plugin that carries credentials (`vikunja`, whose dashboard ships a co-located `.env` file) is enabled in both profiles: plugin credentials, like skill credentials, live inside the tree that ships them, not in a profile-level secret store.

8 The Research Workflow

The designated workflow for math projects is the `math-research-workflow` skill (v1.6), a single-agent adaptation of the MathAgent multi-agent system with Graph-of-Agents message passing and Reflexion quality gates. Its six stages, each gated by a 5-check Reflexion gate:

Stage	Activity	Primary tools
1. Problem decomposition	Restate the problem; open Vikunja tasks and a hypothesis tree (conjecture = parent, attack branches = children with status)	vikunja MCP
2. Literature synthesis	Local-first retrieval cascade: LightRAG → academic-research-hub → Semantic Scholar → SearXNG → ZIMI; every citation passes the <i>Citation Verification Gate</i> (arXiv API / Crossref / S2 before any bibliographic fact is recorded)	lightrag, searxng, zimi, s2 skill
3. Computation	Deterministic-first: SymPy → SageMath → Wolfram Alpha → sandboxed scientific-coding; Error-Feedback Loop with a 3-retry budget, exact prior error carried into each retry	sympy-mcp, sagemath-mcp, wolfram-alpha, scientific-coding
4. Formal verification	Three-stage Lean pipeline: formalize → compile clean (<code>lean4_check</code> against an explicit Mathlib project) → statement-fidelity diff vs. the informal claim; only then is a theorem ledger-marked “proven & Lean-verified”	lean4, lean4-goedel-agent
5. Manuscript production	Drafter → Enhancer → Reviewer pass over LaTeX with the ghasemi-latex-style anti-pattern checklist and notation audit	latex-manuscript MCP, ghasemi-latex-style
6. Reflection & close-out	Claims ledger updated; failures promoted to skill pitfalls or project reports; project README + AGENTS.md status synced in the same commit	—

Supporting state: a per-project failure scratchpad under `~/.hermes/profiles/math/scratchpad/`, per-project verified-source registries for citation dedup, and the cron jobs (weekly arXiv scan in the five-topic overlap set; bi-weekly wiki maintenance) that feed Stage 2 continuously. The Irene package (`~/Code/Python/Irene/`) is the shared numerical backend for polynomial optimization experiments across projects (moment/SOS hierarchies, SONC relaxations), invoked from Irene’s venv rather than through MCP — the intended `hermes-irene` bridge being the broken one documented in Section 6.

9 Identified Couplings and Defects

Ranked by severity, with the evidence for each entry given in the consequence column:

Table 6: Coupling points that bind the profiles to the current project layout.

Coupling	Consequence	Severity
Cross-profile script import (librarian → math)	Every adapter tool script, every native server script and the adapter itself resolve under the math profile; the librarian has no scripts directory of its own	critical
Stale development-worktree path	Both profiles register a tool against a script in a worktree that was merged and removed; dormant in math, failing at startup in the librarian	high
Absolute skill paths in MCP configuration	Every registration hardcodes the full home path; no interpolation mechanism exists	high
Hardcoded service endpoints	Environment files embed host and port directly; no name-based resolution	medium
Project-pinned working directories	Working directory, repository scan roots and execution mode reference specific project trees; the kanban orchestrator is fixed to one profile in both configurations	medium
Toolchain baked into the profile home	The local computation toolchain ships as profile state rather than as a project binding	medium
Stale disabled entries (curator tombstones)	Disabled names that resolve only under the curator archive would silently disable a same-named skill if it were reinstalled	low
Orphaned skill reference	A skill named by the research workflow exists in neither catalog	low
Asymmetric service registration	The same service name means different things in the two profiles	low

10 Toward an Abstraction Layer

The goal is a design in which *profiles are personas and projects are bindings*. A profile declares **what it can do** (capabilities) and *how it behaves* (persona); a project declares **where things live** (paths, endpoints, ledgers). Nothing in today’s configuration should need editing to attach the existing math or librarian profile to a new math project.

A. Capability-gated skill activation (replaces negative lists)

Each `SKILL.md` gains one frontmatter field:

```
requires: [computation, formal-proof] # empty list = always available
```

and each profile declares a positive capability set instead of a disabled list:

Persona	Capabilities
mathematician (today: math)	literature, computation, formal-proof, manuscript, knowledge-mgmt, optimization-backends
librarian	literature, ingestion, knowledge-mgmt, manuscripts-light

The loader activates a skill iff `requires` $\subseteq$ `persona.capabilities`. The librarian's current negative list disappears; tombstone names (entries that resolve only under `.archive/`) become impossible to carry forward, since a capability gate cannot name a skill the persona does not declare; and a new persona is one YAML block, not a diff of every registration.

B. Path interpolation and a shared adapter (removes the critical coupling)

Introduce `${PROFILE_HOME}` and `${HERMES_ROOT}` tokens in `mcp_servers` entries, mirroring the `{baseDir}` mechanism skills already use. Concretely:

1. Move the generic adapter to a shared location: `~/.hermes/tools/mcp_cli_adapter.py` (next to `tirth/uv`, which are already cross-profile). Adapter-wrapped tool scripts move with their skills — they do, since `TOOL_SCRIPT` paths become `${PROFILE_HOME}/skills/...`
2. Each profile's config then references only its own tree; the librarian stops importing from math. Verification is trivial: no path in a profile's config may contain another profile's name.

C. Service registry (removes the endpoint coupling)

A `services.yaml` naming services abstractly, with endpoints resolved from `env` at load time:

```
1 # ~/.hermes/services.yaml
2 lightrag: { url: ${LIGHTRAG_URL}, key: ${LIGHTRAG_API_KEY} }
3 searxng:  { url: ${SEARXNG_URL} }
4 zimi:     { url: ${ZIMI_URL} }
5 vikunja:  { url: ${VIKUNJA_URL}, token: ${VIKUNJA_TOKEN} }
6 siyuan:   { url: ${SIYUAN_URL}, token: ${SIYUAN_TOKEN} }
7 calibre:  { url: ${CALIBRE_BASE_URL} }
8 honcho:   { url: ${HONCHO_BASE_URL} }
```

MCP entries then reference `$services.lightrag.url` etc. The `.env` files remain the secret store (never committed), but IPs move out of configuration logic into environment, so a service migration is an `env` edit, not a per-profile config surgery. This also fixes the calibre asymmetry: one definition, activated by capability ingestion.

D. Toolchain bundle (removes the sandbox coupling)

The sandbox home is a *capability artifact* (computation, formal proof and manuscript), not profile state. Split it: keep identity files (`.gitconfig`, `.ssh`, `.bashrc`) in the profile; extract the toolchain into an installable bundle `math-toolchain.tar` or a Docker/ conda-spec manifest that any persona declaring the relevant capabilities can provision on demand — and that a project without formal-proof needs (e.g. pure numerical work) never pays for.

E. Project binding files (the actual project-independence mechanism)

A new math project is then attached by a single file at its root:

```

1 # ~/Code/<NewProject>/ .agent.yaml
2 profiles: [math, librarian]
3 working_dir: .
4 vikunja_project: <id>           # task ledger target
5 lightRAG_collection: <project>  # literature namespace
6 sage_env: sage                 # toolchain bundle component (conda env name)

```

Profiles discover projects by scanning their `repo_scan_roots`; a profile picks up a project's binding at session start and overrides `terminal.cwd`, ledger targets, and collection names from it. Nothing inside the profiles is edited when the project changes — which is exactly the property today's coupling list (Table 6) shows to be absent.

F. Remediation prerequisites

Before any of A–E is built, four live defects should be fixed in place (each a one-line or one-file change): repoint `hermes-irene` at the merged Irene tree (or disable it for the librarian), prune the tombstone-only names from the disabled lists, remove `simplerag-memory` from `math-research-workflow`'s related-skill list, and give math a real calibre server command or delete its env-only block. These are recorded here as open work items; none is applied by this documentation project.

Instantiation model, end state

Figure 2 summarizes the target relationship between templates, personas and projects.




Figure 2: Target model: personas are instantiated from templates with a capability set; projects bind to personas through root-level binding files. No profile-internal configuration changes when the project changes.

1.1 Roadmap for AgentDesign Itself

Phase	Activity	Status
I. Exploration	Read every configuration, skill tree, state database and script; extract the design data programmatically	done
II. Documentation and packs	This document, the sanitized profile packs of Section 4, the skill inventory and the extracted data	this deliverable
III. Remediation PRs	The four one-line fixes of Section 10(F), as separate commits against the live profiles	open
IV. Abstraction pilot	Implement capability frontmatter + loader gate for one skill pair (e.g. <code>sagemath-mcp</code> under a test persona); verify librarian behavior is unchanged	open
V. Migration	Shared adapter move, <code>services.yaml</code> , toolchain bundle; then retire disabled lists	open

Sources and Verification

Every statement in this document is derived from a live installation of the two profiles rather than from recollection: the configuration files, environment files, profile metadata, scheduled-job definitions, script trees, plugin trees, sandbox homes and complete skill trees of both profiles were read directly, together with the shared Hermes plugin pool and the shared MCP interpreter environment. The extraction is reproducible — the script that produced the comparison data and the skill inventory ships in this project, and the same script rebuilds the two profile packs of Section 4.

- **Extraction** — a script in `references/` reads both profiles and writes the comparison data and the per-entry skill inventory (`skill_inventory.md`).
- **Packs** — the pack builder in `references/` produces the `math/` and `librarian/` folders from the same sources (`build_profile_packs.py`).
- **Structural claims** — the cross-profile path coupling, the dangling registration and the curator tombstones were each confirmed by direct filesystem and configuration inspection, not inferred.

Redaction note: no credential, endpoint, hostname or personal identifier appears anywhere in this document or in the shipped packs. Secrets are replaced by placeholders in the packs and masked in the extracted data artifacts; the redaction pass is applied mechanically by the pack builder and verified by a residual-secret scan over the result.